

classifier

March 13, 2019

```
In [44]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.path as mpath
plt.rcParams['figure.figsize'] = [17, 5]

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.grid_search import GridSearchCV
from sklearn.metrics import precision_recall_curve, average_precision_score
```

```
In [45]: # load data
df_sub1 = pd.read_csv('./labeled_data/sub1label.csv')
df_sub2 = pd.read_csv('./labeled_data/sub1label.csv')

# drop unnecessary columns
df_sub1.drop('Unnamed: 0', axis=1, inplace=True)
df_sub2.drop('Unnamed: 0', axis=1, inplace=True)
```

```
df_sub1.head(10)
```

```
Out[45]:
```

| | Time | A/M | Other | A/M | delta | theta | low_alpha | high_alpha | \ |
|---|----------|-----|-------|------|-------|----------|-----------|------------|---|
| 0 | 4.880859 | | -44.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 1 | 4.882812 | | -38.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 2 | 4.884766 | | -27.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 3 | 4.886719 | | -25.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 4 | 4.888672 | | -19.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 5 | 4.890625 | | -22.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 6 | 4.892578 | | -21.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 7 | 4.894531 | | -8.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 8 | 4.896484 | | 4.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |
| 9 | 4.898438 | | 7.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | |

| | low_beta | high_beta | low_gamma | mid_gamma | blink_stimulation | \ |
|---|----------|-----------|-----------|-----------|-------------------|---|
| 0 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 | |
| 1 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 | |
| 2 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 | |


| | | | | | |
|---|--------|--------|---------|--------|-------|
| 3 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |
| 4 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |
| 5 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |
| 6 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |
| 7 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |
| 8 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |
| 9 | 6806.0 | 9837.0 | 10228.0 | 1916.0 | 849.0 |




| | blink_strength | label |
|---|----------------|-------|
| 0 | 3.783506e-44 | 0 |
| 1 | 4.764415e-43 | 0 |
| 2 | 3.363116e-44 | 0 |
| 3 | 4.203895e-44 | 0 |
| 4 | 0.000000e+00 | 0 |
| 5 | 1.386716e-38 | 0 |
| 6 | 4.049753e-43 | 0 |
| 7 | 1.191104e-43 | 0 |
| 8 | 3.783506e-44 | 0 |
| 9 | 4.764415e-43 | 0 |




```
In [46]: print(df_sub1.shape)
```




```
(43574, 14)
```




```
In [47]: # plot all features vs time
cols = list(df_sub1)
for c in cols:
    if (c != "Time"):
        df_sub1.plot.line(x='Time', y=c)
plt.show()
```











```
In [48]: # remove time from dfs
df_sub1.drop('Time', axis=1, inplace=True)
df_sub2.drop('Time', axis=1, inplace=True)
cols = list(df_sub1)
df_sub1.head()
```

```
Out[48]:
```

| | A/M | Other | A/M | delta | theta | low_alpha | high_alpha | low_beta | \ |
|---|-----|-------|------|-------|----------|-----------|------------|----------|---|
| 0 | | -44.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | 6806.0 | |
| 1 | | -38.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | 6806.0 | |
| 2 | | -27.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | 6806.0 | |
| 3 | | -25.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | 6806.0 | |
| 4 | | -19.0 | 51.0 | 74.0 | 567109.0 | 74006.0 | 38310.0 | 6806.0 | |

| | high_beta | low_gamma | mid_gamma | blink_stimulation | blink_strength | label |
|---|-----------|-----------|-----------|-------------------|----------------|-------|
| 0 | 9837.0 | 10228.0 | 1916.0 | 849.0 | 3.783506e-44 | 0 |
| 1 | 9837.0 | 10228.0 | 1916.0 | 849.0 | 4.764415e-43 | 0 |
| 2 | 9837.0 | 10228.0 | 1916.0 | 849.0 | 3.363116e-44 | 0 |
| 3 | 9837.0 | 10228.0 | 1916.0 | 849.0 | 4.203895e-44 | 0 |
| 4 | 9837.0 | 10228.0 | 1916.0 | 849.0 | 0.000000e+00 | 0 |

```
In [49]: # correlation between features
corr = df_sub1.corr()
corr.style.background_gradient(cmap='coolwarm')

Out[49]: <pandas.io.formats.style.Styler at 0x1a12ea4978>
```

0.1 Random Forest

```
In [50]: # random forest
train, test = train_test_split(df_sub1, test_size=0.2)
X = train.values[:,0:12]
Y = train.values[:,12]

X_test = test.values[:,0:12]
Y_test = test.values[:,12]
clf = RandomForestClassifier(n_estimators=200, max_depth=2, random_state=0)
clf.fit(X, Y)

# feature importance
for importance, feature in zip(clf.feature_importances_, cols):
    print(feature, importance)

A/M Other 0.0
A/M 0.11373261920096323
delta 0.04757553825490368
theta 0.06537066647588381
low_alpha 0.08278091284575494
high_alpha 0.10829649973755964
low_beta 0.08981303492901326
high_beta 0.2022986518637405
low_gamma 0.05447497226715841
mid_gamma 0.13663015040992754
blink_stimulation 0.09902695401509508
blink_strength 0.0
```

```

In [51]: # accuracy
print("mean accuracy: ", clf.score(X_test, Y_test))
Y_scores = []
Y_scores = clf.predict_proba(X_test)[:,-1]
#print(Y_scores)

# cross validation
scores = cross_val_score(clf, X, Y, cv=5)
print("CV score:", scores)

# precision recall
precision, recall, thresholds = precision_recall_curve(Y_test, Y_scores)
average_precision = average_precision_score(Y_test, Y_scores)


plt.step(recall, precision, color='b', alpha=0.2,
         where='post')
plt.fill_between(recall, precision, step='post', alpha=0.2,
                 color='b')

plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('2-class Precision-Recall curve: AP={0:0.2f}'.format(average_precision))
plt.show()

```

mean accuracy: 0.8438324727481354

CV score: [0.80983795 0.85257421 0.82513269 0.82929278 0.83660881]



```

In [52]: # perform grid search for best hyperparams (DO NOT UNCOMMENT UNLESS GRID SEARCH IS NEEDED)
#param_grid = {
#    'n_estimators': [200, 500, 800],
#    'max_features': ['auto', 'sqrt', 'log2'],

```



```

#     'max_depth': [1, 2, 3]
#}
#CV_rfc = GridSearchCV(estimator=clf, param_grid=param_grid, cv=5)
#CV_rfc.fit(X, Y)
#print(CV_rfc.best_params_)

```

```

In [53]: # random forest with grid searched hyperparams
clf = RandomForestClassifier(n_estimators=500, max_depth=3, max_features="auto", random
clf.fit(X, Y)
print("mean accuracy: ", clf.score(X_test, Y_test))

# feature importance
for importance, feature in zip(clf.feature_importances_, cols):
    print(feature, importance)

```

```

mean accuracy: 0.8882386689615606
A/M Other 0.00044221182517868664
A/M 0.10789110554145044
delta 0.07513773687476356
theta 0.0722319514512788
low_alpha 0.08869703499853694
high_alpha 0.09086522306666425
low_beta 0.10028891242252662
high_beta 0.13876202998324882
low_gamma 0.07589378980462765
mid_gamma 0.12011269175824656
blink_stimulation 0.12967731227347776
blink_strength 0.0

```

```

In [54]: # accuracy
print("mean accuracy: ", clf.score(X_test, Y_test))
Y_scores = []
Y_scores = clf.predict_proba(X_test)[:,-1]
#print(Y_scores)

# cross validation
scores = cross_val_score(clf, X, Y, cv=5)
print("CV score:", scores)

# precision recall
precision, recall, thresholds = precision_recall_curve(Y_test, Y_scores)
average_precision = average_precision_score(Y_test, Y_scores)

plt.step(recall, precision, color='b', alpha=0.2,
         where='post')
plt.fill_between(recall, precision, step='post', alpha=0.2,
                 color='b')

```


```

plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('2-class Precision-Recall curve: AP={0:0.2f}'.format(average_precision))
plt.show()

```

mean accuracy: 0.8882386689615606

CV score: [0.89043453 0.87638032 0.87878353 0.88394778 0.88853823]



```

In [55]: # testing on subject 2
train2, test2 = train_test_split(df_sub2, test_size=0.2)
X2 = train2.values[:,0:12]
Y2 = train2.values[:,12]

X_test2 = test2.values[:,0:12]
Y_test2 = test2.values[:,12]

clf.fit(X2, Y2)

# accuracy
print("mean accuracy: ", clf.score(X_test2, Y_test2))
Y_scores2 = []
Y_scores2 = clf.predict_proba(X_test2)[:,-1]
#print(Y_scores)

# cross validation
scores = cross_val_score(clf, X, Y, cv=5)
print("CV score:", scores)

precision, recall, thresholds = precision_recall_curve(Y_test2, Y_scores2)

```

```

average_precision = average_precision_score(Y_test2, Y_scores2)

plt.step(recall, precision, color='b', alpha=0.2,
         where='post')
plt.fill_between(recall, precision, step='post', alpha=0.2,
                 color='b')

plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title('2-class Precision-Recall curve: AP={0:0.2f}'.format(average_precision))
plt.show()

```

mean accuracy: 0.8772231784279977

CV score: [0.89043453 0.87638032 0.87878353 0.88394778 0.88853823]

